

Chapter 6: The basic SQL

6.1 SQL Data Definition and Data Types

SQL uses the terms table, row and column for the formal relational model terms *relation*, *tuple*, and *attribute*.

The main SQL command for data definition is the **CREATE** statement, which can be used to create schemas, tables (relations), types, and domains, as well as other constructs such as views, assertion and triggers.

6.1.1 Schema and Catalog Concepts in SQL

The concept of an SQL schema was incorporated starting with SQL2 in order to **group together tables** and **other constructs** that belong to the **same database application** (in some system, a *schema* is called a *database*).

An **SQL schema** is identified by a **schema name** and includes an **authorization identifier** to indicate the user or account who owns the schema, as well as **descriptors** for *each element* in the schema.

Schema elements include *tables, types, constraints, views, domains, and other constructs* (such as authorization grants) that describe the schema.

Can be created via the **CREATE SCHEMA**, element can be defined later.

```
CREATE SCHEMA COMPANY AUTHORIZATION 'Jsmith';
```

NOTE: In general, not all user are authorized to create schemas, must be explicitly granted the relevant user account by the DBA or system administrator.

SQL use the concept of a **catalog** - a named collection of schemas.

A **catalog** always contains a *special schema* called **INFORMATION_SCHEMA** - which provides information on all the schemas in the catalog and all the element descriptors in these schema.

Integrity constraints such as **referential integrity** can be defined between relations only if they exist in schema with **same catalog**. And schema with the same catalog can share certain elements.

6.1.2 The CREATE TABLE Command in SQL

The **CREATE TABLE** command is used to specify a new relation by giving it a name and specifying its attributes and initial constraints.

The **attributes** are **specified first** and *each attribute is given a name, a data type to specify its domain of values and possibly attribute constraints, such as **NOT NULL**.*

The *key, entity integrity, and referential integrity constraints* can be specified within the **CREATE TABLE** statement after the attributes are declared, or they can be *added later* using the **ALTER TABLE** command.

The **SQL schema** in which the relations are declared is implicitly specified in the environment in which the **CREATE TABLE** statements are executed. Alternative, we can *explicitly attach the schema name to the relation name, separated by a period*.

```
CREATE TABLE COMPANY.EMPLOYEE  
rather than  
CREATE TABLE EMPLOYEE
```

```

CREATE TABLE EMPLOYEE
( Fname          VARCHAR(15)          NOT NULL,
  Minit          CHAR,
  Lname          VARCHAR(15)          NOT NULL,
  Ssn            CHAR(9)              NOT NULL,
  Bdate          DATE,
  Address        VARCHAR(30),
  Sex            CHAR,
  Salary         DECIMAL(10,2),
  Super_ssn      CHAR(9),
  Dno            INT                  NOT NULL,
  PRIMARY KEY (Ssn),
CREATE TABLE DEPARTMENT
( Dname          VARCHAR(15)          NOT NULL,
  Dnumber        INT                  NOT NULL,
  Mgr_ssn        CHAR(9)              NOT NULL,
  Mgr_start_date DATE,
  PRIMARY KEY (Dnumber),
  UNIQUE (Dname),
  FOREIGN KEY (Mgr_ssn) REFERENCES EMPLOYEE(Ssn) );
CREATE TABLE DEPT_LOCATIONS
( Dnumber        INT                  NOT NULL,
  Dlocation      VARCHAR(15)          NOT NULL,
  PRIMARY KEY (Dnumber, Dlocation),
  FOREIGN KEY (Dnumber) REFERENCES DEPARTMENT(Dnumber) );
CREATE TABLE PROJECT
( Pname          VARCHAR(15)          NOT NULL,
  Pnumber        INT                  NOT NULL,
  Plocation      VARCHAR(15),
  Dnum           INT                  NOT NULL,
  PRIMARY KEY (Pnumber),
  UNIQUE (Pname),
  FOREIGN KEY (Dnum) REFERENCES DEPARTMENT(Dnumber) );
CREATE TABLE WORKS_ON
( Essn           CHAR(9)              NOT NULL,
  Pno            INT                  NOT NULL,
  Hours          DECIMAL(3,1)         NOT NULL,
  PRIMARY KEY (Essn, Pno),
  FOREIGN KEY (Essn) REFERENCES EMPLOYEE(Ssn),
  FOREIGN KEY (Pno) REFERENCES PROJECT(Pnumber) );
CREATE TABLE DEPENDENT
( Essn           CHAR(9)              NOT NULL,
  Dependent_name VARCHAR(15)          NOT NULL,
  Sex            CHAR,
  Bdate          DATE,
  Relationship    VARCHAR(8),
  PRIMARY KEY (Essn, Dependent_name),
  FOREIGN KEY (Essn) REFERENCES EMPLOYEE(Ssn) );

```

Figure 6.1
SQL CREATE
TABLE data
definition statements
for defining the
COMPANY schema
from Figure 5.7.

The relations declared through **CREATE TABLE** statements are called **base tables** (base relations) - means that **the table** and **its rows** are actually created and

stored as a file by the DMBS.

Base relations are distinguished from **virtual relations**, create through the **CREATE VIEW** statement - which may or may not correspond to an actual physical file.

In **SQL**, the *attributes* in a **base tables** are considered to be *ordered in the sequence in which they are specified* in the **CREATE TABLE** statement. However, **rows (tuples)** are **not** considered to be *ordered in the sequence*.

6.1.3 Attribute Data Types and Domains in SQL

The **basic data types** available for attributes include numeric, character string, bit string, Boolean, date, and time:

- **Numeric** data types include integer numbers of various sizes (**INTEGER** or **INT**, and **SMALLINT**) and floating-point (real) numbers of various precision (**FLOAT** or **REAL**, and **DOUBLE PRECISION**). **Formatted numbers** can be declared by using **DECIMAL(i, j)**—or **DEC(i, j)** or **NUMERIC(i, j)**—where *i*, **the precision, is the total number of decimal digits** and *j*, **the scale, is the number of digits after the decimal point**. The **default for scale is zero**, and the **default for precision is implementation-defined**.
- **Character - string** data types are either fixed length—**CHAR(n)** or **CHARACTER(n)**, where **n is the number of characters**—or varying length—**VARCHAR(n)** or **CHAR VARYING(n)** or **CHARACTER VARYING(n)**, where **n is the maximum number of characters**. When specifying a literal string value, it is placed between **single quotation marks (apostrophes)**, and it is **case sensitive** (a distinction is made **between uppercase and lowercase**). For fixed length strings, **a shorter string is padded with blank characters to the right**. For example, if the value 'Smith' is for an attribute of type **CHAR(10)**, it is padded with five blank characters to become 'Smith' if needed. **Padded blanks are generally ignored when strings are compared**. For comparison purposes, strings are considered **ordered in alphabetic (or lexicographic) order**; if a string str1 appears before another string str2 in alphabetic order, then str1 is considered to be less than str2. There is **also a concatenation operator denoted by || (double vertical bar) that can concatenate two strings**. For example, **'abc' || 'XYZ' results in a single string 'abcXYZ'**. Another variable-length string data type called **CHARACTER LARGE OBJECT**

or **CLOB** is also available to specify columns that have large text values, such as documents. The **CLOB** maximum length can be specified in **kilobytes (K)**, **megabytes (M)**, or **gigabytes (G)**. For example, **CLOB(20M)** specifies a **maximum length of 20 megabytes**.

- **Bit - string** data types are either of **fixed length n—BIT(n)**—or **varying length — BIT VARYING(n)**, where **n** is the maximum number of bits. The default for **n**, the length of a character string or bit string, is 1. Literal bit strings are placed between single quotes but preceded by a **B** to distinguish them from character strings; for example, **B'10101'**. Another variable-length bit string data type called **BINARY LARGE OBJECT** or **BLOB** is also available to specify columns that have **large binary values**, such as images. As for **CLOB**, the maximum length of a **BLOB** can be specified in **kilobits (K)**, **megabits (M)**, or **gigabits (G)**. For example, **BLOB(30G)** specifies a maximum length of 30 gigabits.
- **Boolean** data type has the traditional values of **TRUE** or **FALSE**. In SQL, because of the presence of **NULL values**, a three-valued logic is used, so a **third possible value for a Boolean data type is UNKNOWN**.
- **Date** data type has **ten positions**, and its components are **YEAR, MONTH, and DAY** in the form **YYYY-MM-DD**. The **TIME** data type has **at least eight positions**, with the components **HOUR, MINUTE, and SECOND** in the form **HH:MM:SS**. Only valid dates and times should be allowed by the SQL implementation. This implies that **months should be between 1 and 12** and **days must be between 01 and 31**; furthermore, a day should be a valid day for the corresponding month. The **< (less than) comparison can be used with dates or times**—an earlier date is considered to be smaller than a later date, and similarly with time. Literal values are represented by **single quoted strings preceded by the keyword DATE or TIME**; for example, **DATE '2014-09-27'** or **TIME '09:12:47'**. In addition, a data type **TIME(i)**, where *i* is called *time fractional seconds precision*, specifies *i + 1* additional positions for **TIME**—one position for an additional period (.) separator character, and *i* positions for specifying decimal fractions of a second. A **TIME WITH TIME ZONE** data type includes an additional **six positions for specifying the displacement from the standard universal time zone**, which is in the range **+13:00 to -12:59** in units

of **HOURS:MINUTES**. If **WITH TIME ZONE** is not included, the default is the local time zone for the SQL session.

- **Timestamp** data type (**TIMESTAMP**) include the **DATE** and **TIME** fields, plus a minimum of six positions for the decimal fraction of seconds and an optional **WITH TIME ZONE**. For example, **TIMESTAMP '2019-09-27 09:12:47.648302'**.
- Another data type related to **DATE**, **TIME**, and **TIMESTAMP** is the **INTERVAL** data type. This specifies an **interval**—*a relative value that can be used to increment or decrement an absolute value of a date, time, or timestamp*. **Intervals** are qualified to be either **YEAR/MONTH** intervals or **DAY/TIME** intervals

The format of **DATE**, **TIME**, and **TIMESTAMP** can be considered as a **special type of string**. Hence, they can generally be used in **string comparisons by being cast (or coerced or converted)** into the equivalent strings

It is possible to specify the data type of each attribute directly, as in Figure 6.1; alternatively, a **domain can be declared, and the domain name can be used with the attribute specification**. This makes it **easier to change the data type** for a domain that is used by **numerous attributes** in a schema, and improves schema readability. For example, we can create a domain **SSN_TYPE** by the following statement.

```
CREATE DOMAIN SSN_TYPE AS CHAR(9);
```

In **SQL**, there is also a **CREATE TYPE** command, which can be used to create user defined types or UDTs. These can then be used either as **data types for attributes**, or as **the basis for creating tables**.

6.2 Specifying Constraints in SQL

6.2.1 Specifying Attribute Constraints and Attribute Defaults

SQL allows **NULLs** as attribute values, a *constraint* **NOT NULL** may be specified if **NULL** is not permitted for a particular attribute. Always implicitly specified for the attributes that are part of the *primary key*, but it can be specified for any attributes whose values are *required not to be NULL*.

It is also possible to **define a default value** for an attribute by appending the clause **DEFAULT <value> to an attribute definition**. The **default value** is included in any new tuple if an explicit value is not provided for that attribute. If **no default clause is specified**, the default *default* value is **NULL for attributes that do not have the NOT NULL constraint**.

```
CREATE TABLE EMPLOYEE
( ... ,
  Dno          INT          NOT NULL      DEFAULT 1,
  CONSTRAINT EMPPK
    PRIMARY KEY (Ssn),
  CONSTRAINT EMPSUPERFK
    FOREIGN KEY (Super_ssn) REFERENCES EMPLOYEE(Ssn)
      ON DELETE SET NULL      ON UPDATE CASCADE,
  CONSTRAINT EMPDEPTFK
    FOREIGN KEY (Dno) REFERENCES DEPARTMENT(Dnumber)
      ON DELETE SET DEFAULT   ON UPDATE CASCADE);
CREATE TABLE DEPARTMENT
( ... ,
  Mgr_ssn CHAR(9)          NOT NULL      DEFAULT '888665555',
  ... ,
  CONSTRAINT DEPTPK
    PRIMARY KEY (Dnumber),
  CONSTRAINT DEPTSK
    UNIQUE (Dname),
  CONSTRAINT DEPTMGRFK
    FOREIGN KEY (Mgr_ssn) REFERENCES EMPLOYEE(Ssn)
      ON DELETE SET DEFAULT   ON UPDATE CASCADE);
CREATE TABLE DEPT_LOCATIONS
( ... ,
  PRIMARY KEY (Dnumber, Dlocation),
  FOREIGN KEY (Dnumber) REFERENCES DEPARTMENT(Dnumber)
    ON DELETE CASCADE        ON UPDATE CASCADE);
```

Figure 6.2
Example illustrating how default attribute values and referential integrity triggered actions are specified in SQL.

Another type of constraint can restrict attribute or domain values using the CHECK clause following an attribute or domain definition. For example, suppose that department numbers are restricted to integer numbers between 1 and 20; then, we can change the attribute declaration of Dnumber in the DEPARTMENT table (see Figure 6.1) to the following:

```
Dnumber INT NOT NULL CHECK (Dnumber > 0 AND Dnumber < 21);
```

The **CHECK** clause can also be used in **conjunction** with the **CREATE DOMAIN** statement. For example, we can write the following statement:

```
CREATE DOMAIN D_NUM AS INTEGER
CHECK (D_NUM > 0 AND D_NUM < 21);
```

6.2.2 Specifying Key and Referential Integrity Constraints

The **PRIMARY KEY** clause specifies **one or more attributes that make up the primary key** of a relation. If a **primary key has a single attribute**, the clause can follow the attribute directly.

```
Dnumber INT PRIMARY KEY,
```

The **UNIQUE** clause specifies alternate (unique) keys, also known as candidate keys. The **UNIQUE** clause can also be specified directly for a **unique key if it is a single attribute**, as in the following example:

```
Dname VARCHAR(15) UNIQUE,
```

Referential integrity is specified via the **FOREIGN KEY** clause.

Referential integrity is specified via the **FOREIGN KEY** clause, as shown in Figure 6.1. As we discussed in Section 5.2.4, a **referential integrity constraint** can be **violated when tuples are inserted or deleted**, or when a **foreign key or primary key attribute value is updated**. The default action that **SQL** takes for an **integrity violation** is to **reject the update operation** that will cause a violation, which is known as the **RESTRICT** option. However, the schema designer can specify an alternative action to be taken by attaching a referential triggered action clause to any foreign key constraint. The options include **SET NULL**, **CASCADE**, and **SET DEFAULT**. An option must be qualified with either **ON DELETE** or **ON UPDATE**. We illustrate this with the examples shown in Figure 6.2. Here, the database designer chooses **ON DELETE SET NULL** and **ON UPDATE CASCADE** for the foreign key **Super_ssn** of **EMPLOYEE**. This means that if the tuple for a supervising employee is **deleted**, the value of **Super_ssn** is **automatically set to NULL for all employee tuples that were referencing the deleted employee tuple**. On the other hand, if the **Ssn** value for a supervising employee is **updated** (say, because it was entered incorrectly), the **new value is cascaded to Super_ssn for all employee tuples referencing the updated employee tuple**.

6.2.3 Giving Names to Constraints

Figure 6.2 also illustrates how a constraint may be given a constraint name, following the keyword **CONSTRAINT**. The **names of all constraints within a particular schema must be unique**. A constraint name is used to **identify a particular constraint** in case the constraint **must be dropped later and replaced with another constraint**. Giving names to constraints is **optional**. It is also possible to temporarily defer a constraint until the end of a transaction

6.2.4 Specifying Constraints on Tuples Using CHECK

In addition to key and referential integrity constraints, which are specified by special keywords, other **table constraints** can be specified through additional **CHECK** clauses **at the end of a CREATE TABLE statement**. These can be called **row-based constraints** because *they apply to each row individually and are checked whenever a row is inserted or modified*. For example, suppose that the **DEPARTMENT** table in Figure 6.1 had an additional attribute `Dept_create_date`, which stores the date when the department was created. Then we could add the following **CHECK** clause at the end of the **CREATE TABLE** statement for the **DEPARTMENT** table to make sure that a manager's start date is later than the department creation date.

6.3 Basic Retrieval Queries in SQL

SQL has **one basic statement** for retrieving information from a database: the **SELECT** statement. The **SELECT statement is not the same as the SELECT operation of relational algebra**.

Before proceeding, we must **point out an important distinction between the practical SQL model and the formal relational model**. **SQL allows a table (relation) to have two or more tuples that are identical in all their attribute values**. Hence, in general, an **SQL table is not a set of tuples, because a set does not allow two identical members**; rather, it is a **multiset** (sometimes called a **bag**) of tuples. Some **SQL relations are constrained to be sets** because a **key constraint has been declared** or because the **DISTINCT** option has been used with the **SELECT** statement.

6.3.1 The SELECT-FROM-WHERE Structure of Basic SQL Queries

The basic form of the **SELECT** statement, sometimes called a **mapping** or a **select-from-where** block, is formed of the three clauses **SELECT**, **FROM**, and **WHERE** and has the following form:

```
SELECT    <attribute list>
FROM      <table list>
WHERE     <condition>;
```

where

- <attribute list> is a list of attribute names whose values are to be retrieved by the query.
- <table list> is a list of the relation names required to process the query.
- <condition> is a conditional (Boolean) expression that identifies the tuples to be retrieved by the query.

Query 0. Retrieve the birth date and address of the employee(s) whose name is 'John B. Smith'.

```
Q0:    SELECT    Bdate, Address
        FROM      EMPLOYEE
        WHERE     Fname = 'John' AND Minit = 'B' AND Lname = 'Smith';
```

The **SELECT** clause of **SQL** specifies the attributes whose values are to be retrieved, which are called **the projection attributes** in relational algebra

WHERE clause specifies the Boolean condition that must be true for any retrieved tuple, which is known as **the selection condition** in relational algebra

Query 1. Retrieve the name and address of all employees who work for the 'Research' department.

```
Q1:    SELECT    Fname, Lname, Address
        FROM      EMPLOYEE, DEPARTMENT
        WHERE     Dname = 'Research' AND Dnumber = Dno;
```

In the **WHERE** clause of Q1, the condition $Dname = \text{'Research'}$ is a **selection condition** that chooses the particular tuple of interest in the DEPARTMENT table, because Dname is an attribute of DEPARTMENT. The condition $Dnumber = Dno$ is called a **join condition**, because it combines two tuples: one from DEPARTMENT

and one from EMPLOYEE, whenever the value of Dnumber in DEPARTMENT is equal to the value of Dno in EMPLOYEE.

A query that involves **only selection and join conditions plus projection attributes** is

known as a

select-project-join query. The next example is a select-project-join query with two join conditions.

Query 2. For every project located in 'Stafford', list the project number, the controlling department number, and the department manager's last name, address, and birth date.

```
Q2:  SELECT  Pnumber, Dnum, Lname, Address, Bdate
      FROM    PROJECT, DEPARTMENT, EMPLOYEE
      WHERE   Dnum = Dnumber AND Mgr_ssn = Ssn AND
             Plocation = 'Stafford'
```


In **SQL**, the **same name can be used** for two (or more) attributes as long as the attributes are in **different tables**. If this is the case, and a multitable query refers to two or more attributes with the same name, we *must* **qualify** the attribute name **with the relation name to prevent ambiguity**. This is done by **prefixing the relation name to the attribute name and separating the two by a period**.

```
Q1A:  SELECT  Fname, EMPLOYEE.Name, Address
        FROM    EMPLOYEE, DEPARTMENT
        WHERE   DEPARTMENT.Name = 'Research' AND
                DEPARTMENT.Dnumber = EMPLOYEE.Dnumber;
```

Fully qualified attribute names can be **used** for clarity even **if there is no ambiguity** in attribute names. Q1 can be rewritten as Q1' below with fully qualified attribute names. We can **also rename the table names to shorter names by creating an alias for each table name** to avoid repeated typing of long table names

```
Q1':   SELECT  EMPLOYEE.Fname, EMPLOYEE.LName,
                EMPLOYEE.Address
        FROM    EMPLOYEE, DEPARTMENT
        WHERE   DEPARTMENT.DName = 'Research' AND
                DEPARTMENT.Dnumber = EMPLOYEE.Dno;
```

The ambiguity of attribute names also arises in the case of queries that refer to the same relation twice, as in the following example

Query 8. For each employee, retrieve the employee's first and last name and the first and last name of his or her immediate supervisor.

```
Q8:    SELECT  E.Fname, E.Lname, S.Fname, S.Lname
        FROM    EMPLOYEE AS E, EMPLOYEE AS S
        WHERE   E.Super_ssn = S.Ssn;
```

In this case, we are required to declare alternative relation names E and S, called aliases or tuple variables, for the EMPLOYEE relation. An **alias** can follow the keyword **AS**, as shown in Q8, or it can **directly follow the relation name**—for

example, by writing **EMPLOYEE E**, **EMPLOYEE S** in the **FROM** clause of Q8. It is **also possible to rename the relation attributes** within the query in **SQL** by giving them **aliases**. For example, if we write

```
EMPLOYEE AS E(Fn, Mi, Ln, Ssn, Bd, Addr, Sex, Sal, Sssn, Dno)
```

in the **FROM** clause, **Fn** becomes an alias for **Fname**, **Mi** for **Minit**, **Ln** for **Lname**, and so on.

We can use this **alias-naming** or **renaming** mechanism in any **SQL** query to specify tuple variables for every table in the **WHERE** clause, whether or not the same relation needs to be referenced more than once.

```
Q1B:  SELECT  E.Fname, E.LName, E.Address
      FROM    EMPLOYEE AS E, DEPARTMENT AS D
      WHERE   D.DName = 'Research' AND D.Dnumber = E.Dno;
```

A missing **WHERE** clause indicates **no condition on tuple selection**; hence, all tuples of the relation specified in the **FROM** clause **qualify and are selected for the query result**. If more than one relation is **specified** in the **FROM** clause and **there is no WHERE clause**, then the **CROSS PRODUCT**—all possible tuple combinations—of these relations is selected.

Queries 9 and 10. Select all **EMPLOYEE Ssns** (Q9) and all combinations of **EMPLOYEE Ssn** and **DEPARTMENT Dname** (Q10) in the database.

```
Q9:   SELECT  Ssn
      FROM    EMPLOYEE;

Q10:  SELECT  Ssn, Dname
      FROM    EMPLOYEE, DEPARTMENT;
```

It is extremely important to specify every selection and join condition in the **WHERE** clause; if any such condition is **overlooked, incorrect and very large relations** may result.

To **retrieve all the attribute values of the selected tuples**, we do not have to list the attribute names explicitly in SQL; we just specify an **asterisk (*)**, which stands for all the attributes. *The * can also be prefixed by the relation name or alias; for example, **EMPLOYEE.*** refers to all attributes of the EMPLOYEE table.*

Query

Q1C retrieves all the attribute values of any **EMPLOYEE** who works in **DEPARTMENT number 5** (Figure 6.3(g)), query **Q1D** retrieves all the attributes of an **EMPLOYEE** and the attributes of the **DEPARTMENT** in which he or she works for every employee of the '**Research**' department, and Q10A specifies the **CROSS PRODUCT** of the **EMPLOYEE** and **DEPARTMENT** relations.

```
Q1C:  SELECT  *
      FROM    EMPLOYEE
      WHERE   Dno = 5;

Q1D:  SELECT  *
      FROM    EMPLOYEE, DEPARTMENT
      WHERE   Dname = 'Research' AND Dno = Dnumber;

Q10A:  SELECT  *
      FROM    EMPLOYEE, DEPARTMENT;
```

6.3.4 Tables as Sets in SQL

As we mentioned earlier, **SQL usually treats a table not as a set but rather as a multiset**; *duplicate tuples can appear more than once in a table*, and in the result of a query. SQL **does not automatically eliminate duplicate tuples** in the results of queries, for the following reasons:

- Duplicate elimination is an expensive operation. One way to implement it is to sort the tuples first and then eliminate duplicates.
- The user may want to see duplicate tuples in the result of a query.
- When an aggregate function (see Section 7.1.7) is applied to tuples, in most cases we do not want to eliminate duplicates.

Figure 6.4

Results of additional SQL queries when applied to the COMPANY database state shown in Figure 5.6. (a) Q11. (b) Q11A. (c) Q16. (d) Q18.

(a)	<table><tr><th>Salary</th></tr><tr><td>30000</td></tr><tr><td>40000</td></tr><tr><td>25000</td></tr><tr><td>43000</td></tr><tr><td>38000</td></tr><tr><td>25000</td></tr><tr><td>25000</td></tr><tr><td>55000</td></tr></table>	Salary	30000	40000	25000	43000	38000	25000	25000	55000	(b)	<table><tr><th>Salary</th></tr><tr><td>30000</td></tr><tr><td>40000</td></tr><tr><td>25000</td></tr><tr><td>43000</td></tr><tr><td>38000</td></tr><tr><td>55000</td></tr></table>	Salary	30000	40000	25000	43000	38000	55000	(c)	<table><tr><th>Fname</th><th>Lname</th></tr><tr><td></td><td></td></tr></table>	Fname	Lname		
Salary																									
30000																									
40000																									
25000																									
43000																									
38000																									
25000																									
25000																									
55000																									
Salary																									
30000																									
40000																									
25000																									
43000																									
38000																									
55000																									
Fname	Lname																								
		(d)	<table><tr><th>Fname</th><th>Lname</th></tr><tr><td>James</td><td>Borg</td></tr></table>	Fname	Lname	James	Borg																		
Fname	Lname																								
James	Borg																								

An **SQL** table with a key is restricted to being a set, since the key value must be distinct in each tuple. If we **do want to eliminate duplicate tuples** from the result of an **SQL** query, we use the keyword **DISTINCT** in the **SELECT** clause, meaning that only distinct tuples should remain in the result. In general, a query with **SELECT DISTINCT** eliminates duplicates, whereas a query with **SELECT ALL** does not. Specifying **SELECT** with neither **ALL** nor **DISTINCT**—is equivalent to **SELECT ALL**.

SQL has **directly incorporated** some of the **set operations from mathematical set theory**, which are also part of relational algebra (see Chapter 8). There are **set union (UNION)**, **set difference (EXCEPT)**, and **set intersection (INTERSECT)** operations. The relations resulting from these set operations are **sets of tuples**; that is, **duplicate tuples are eliminated from the result**. These set operations **apply only to type compatible relations**, so we must **make sure that the two relations on which we apply the operation have the same attributes** and that the **attributes appear in the same order in both relations**. The next example illustrates the use of UNION

(a)	R	S	(b)	T	(c)	T
	A	A		A		A
	a1	a1		a1		a2
	a2	a2		a1		a3
	a2	a4		a2		
	a3	a5		a2		
				a2		
				a3	(d)	T
				a4		A
				a5		a1
						a2

Figure 6.5

The results of SQL multiset operations. (a) Two tables, R(A) and S(A). (b) R(A) UNION ALL S(A). (c) R(A) EXCEPT ALL S(A). (d) R(A) INTERSECT ALL S(A).

Query 4. Make a list of all project numbers for projects that involve an employee whose last name is 'Smith', either as a worker or as a manager of the department that controls the project.

```
Q4A:  ( SELECT  DISTINCT Pnumber
      FROM      PROJECT, DEPARTMENT, EMPLOYEE
      WHERE     Dnum = Dnumber AND Mgr_ssn = Ssn
            AND  Lname = 'Smith' )

      UNION

      ( SELECT  DISTINCT Pnumber
      FROM      PROJECT, WORKS_ON, EMPLOYEE
      WHERE     Pnumber = Pno AND Essn = Ssn
            AND  Lname = 'Smith' );
```

The first **SELECT** query retrieves the projects that involve a 'Smith' as manager of the department that controls the project, and the second retrieves the projects that involve a 'Smith' as a worker on the project. Notice that if several employees have the last name 'Smith', the project names involving any of them will be retrieved. Applying the **UNION** operation to the two SELECT queries gives the desired result.

SQL also has corresponding multiset operations, which are followed by the keyword **ALL** (**UNION ALL**, **EXCEPT ALL**, **INTERSECT ALL**). Their **results are multisets (duplicates are not eliminated)**. The behavior of these operations is illustrated by the examples in Figure 6.5. Basically, each tuple—whether it is a duplicate or not—is considered as a different tuple when applying these operations.

Comparison conditions on only parts of a character string, using the **LIKE** comparison operator. This can be used for **string pattern matching**. **Partial strings** are specified **using two reserved characters**: **%** replaces an **arbitrary number of zero or more characters**, and the **underscore (_)** replaces a **single character**. For example, consider the following query.

```
Q12:  SELECT  Fname, Lname
      FROM    EMPLOYEE
      WHERE   Address LIKE '%Houston,TX%';
```

Query 12A. Find all employees who were born during the 1950s.

```
Q12:  SELECT  Fname, Lname
        FROM    EMPLOYEE
        WHERE   Bdate LIKE '____7____';
```

The standard arithmetic operators for **addition (+)**, **subtraction (-)**, **multiplication (*)**, and **division (/)** can be applied to numeric values or attributes with numeric domains. For example, suppose that we want to see the **effect of giving all employees who work on the 'ProductX' project a 10% raise**; we can issue Query

13 to see what their salaries would become. This example also shows how we can rename an attribute in the query result using AS in the SELECT clause.

Query 13. Show the resulting salaries if every employee working on the 'ProductX' project is given a 10% raise.

```
Q13:  SELECT    E.Fname, E.Lname, 1.1 * E.Salary AS Increased_sal
      FROM      EMPLOYEE AS E, WORKS_ON AS W, PROJECT AS P
      WHERE     E.Ssn = W.Essn AND W.Pno = P.Pnumber AND
               P.Pname = 'ProductX';
```

For **string data types**, the **concatenate operator ||** can be used in a query to **append two string values**. For **date, time, timestamp, and interval data types**, operators include **incrementing (+) or decrementing (-)** a date, time, or timestamp by an interval. In addition, **an interval value** is the result of the **difference between two date, time, or timestamp values**. Another comparison operator, which can be used for convenience, is **BETWEEN**, which is illustrated in Query 14

Query 14. Retrieve all employees in department 5 whose salary is between \$30,000 and \$40,000.

```
Q14:  SELECT    *
      FROM      EMPLOYEE
      WHERE     (Salary BETWEEN 30000 AND 40000) AND Dno = 5;
```

The condition (Salary **BETWEEN** 30000 **AND** 40000) in Q14 is equivalent to the condition ((Salary >= 30000) **AND** (Salary <= 40000)).

6.3.6 Ordering of Query Result

SQL allows the user to order the tuples in the result of a query by the values of one or more of the attributes that appear in the query result, by using the **ORDER BY** clause. This is illustrated by Query 15

Query 15. Retrieve a list of employees and the projects they are working on, ordered by department and, within each department, ordered alphabetically by last name, then first name.

```
Q15:  SELECT    D.Dname, E.Lname, E.Fname, P.Pname
      FROM      DEPARTMENT AS D, EMPLOYEE AS E, WORKS_ON AS W,
      PROJECT AS P
      WHERE     D.Dnumber = E.Dno AND E.Ssn = W.Essn AND W.Pno =
      P.Pnumber
      ORDER BY  D.Dname, E.Lname, E.Fname;
```

The **default order** is in **ascending** order of values. We can specify the keyword **DESC** if we want to see the result in a **descending** order of values. The keyword **ASC** can be used to specify **ascending** order explicitly. For example, if we want descending alphabetical order on Dname and ascending order on Lname, Fname, the **ORDER BY** clause of Q15 can be written as

```
ORDER BY D.Dname DESC, E.Lname ASC, E.Fname ASC
```

6.3.7 Discussion and Summary of Basic SQL Retrieval Queries

A *simple* retrieval query in SQL can consist of up to four clauses, but only the first two—SELECT and FROM—are mandatory. The clauses are specified in the following order, with the clauses between square brackets [...] being optional:

```
SELECT    <attribute list>
FROM      <table list>
[ WHERE   <condition> ]
[ ORDER BY <attribute list> ];
```

The SELECT clause lists the attributes to be retrieved, and the FROM clause specifies all relations (tables) needed in the simple query. The WHERE clause identifies the conditions for selecting the tuples from these relations, including

join conditions if needed. **ORDER BY** specifies an order for displaying the results of a query.

6.4 INSERT, DELETE and UPDATE Statements in SQL

INSERT is used to **add a single tuple (row) to a relation (table)**. We must **specify the relation name and a list of values for the tuple**. The values **should be listed in the same order** in which the corresponding attributes were specified in the **CREATE TABLE** command.

```
U1:  INSERT INTO  EMPLOYEE
      VALUES      ('Richard', 'K', 'Marini', '653298653', '1962-12-30', '98
                    Oak Forest, Katy, TX', 'M', 37000, '653298653', 4 );
```

A second form of the **INSERT** statement allows the user to **specify explicit attribute names** that correspond to the values provided in the **INSERT** command. This is **useful if a relation has many attributes** but **only a few of those attributes are assigned values** in the new tuple. However, the values **must include all attributes with NOT NULL specification and no default value**. Attributes with **NULL allowed or DEFAULT values are the ones that can be left out**.

```
U1A: INSERT INTO  EMPLOYEE (Fname, Lname, Dno, Ssn)
      VALUES      ('Richard', 'Marini', 4, '653298653');
```

Attributes **not specified** in U1A are set to their **DEFAULT** or to **NULL**, and the values are listed in the **same order as the attributes are listed** in the **INSERT** command itself. It is also possible to insert into a relation multiple tuples separated by commas in a single **INSERT** command. The attribute values forming each tuple are **enclosed in parentheses**.

```
INSERT INTO EMPLOYEE (Ssn, Fname, Lname, Dno, Salary)
VALUES
  ('111-22-3333', 'Alice', 'Johnson', 2, 55000),
  ('444-55-6666', 'Bob', 'Williams', 4, 60000),
  ('777-88-9999', 'Charlie', 'Brown', 1, 65000);
```

A **DBMS** that **fully implements SQL** should **support and enforce all the integrity constraints** that can be specified in the **DDL**. For example, if we issue the command in U2 on the database shown in Figure 5.6, the **DBMS should reject the operation because no DEPARTMENT tuple exists in the database with Dnumber**

= 2. Similarly, **U2A** would be rejected because no **Ssn** value is provided and it is the **primary key**, which cannot be NULL

```
U2:  INSERT INTO      EMPLOYEE (Fname, Lname, Ssn, Dno)
      VALUES          ('Robert', 'Hatcher', '980760540', 2);
      (U2 is rejected if referential integrity checking is provided by DBMS.)

U2A: INSERT INTO      EMPLOYEE (Fname, Lname, Dno)
      VALUES          ('Robert', 'Hatcher', 5);
      (U2A is rejected if NOT NULL checking is provided by DBMS.)
```

A **variation of the INSERT command** inserts **multiple tuples** into a relation in conjunction with **creating the relation and loading it with the result of a query**.

```
U3A:  CREATE TABLE    WORKS_ON_INFO
      ( Emp_name        VARCHAR(15),
        Proj_name       VARCHAR(15),
        Hours_per_week  DECIMAL(3,1) );

U3B:  INSERT INTO      WORKS_ON_INFO ( Emp_name, Proj_name,
      Hours_per_week )
      SELECT            E.Lname, P.Pname, W.Hours
      FROM              PROJECT P, WORKS_ON W, EMPLOYEE E
      WHERE             P.Pnumber = W.Pno AND W.Essn = E.Ssn;
```

Most **DBMSs** have **bulk loading tools** that allow a user to load formatted data from a file into a table without having to write a large number of **INSERT commands**.

Another variation for loading data is to create a new table **TNEW** that has the same attributes as an **existing table T**, and load some of the data currently in **T** into **TNEW**.

```
CREATE TABLE          D5EMPS LIKE EMPLOYEE
(SELECT
FROM
WHERE                  E.Dno = 5) WITH DATA;
```

The clause **WITH DATA** specifies that **the table will be created and loaded with the data specified** in the query, although in some implementations **it may be left out**.

6.4.2 The DELETE Command

The **DELETE** command removes tuples from a relation. It includes a **WHERE** clause, similar to that used in an **SQL** query, to select the tuples to be deleted. **Tuples are explicitly deleted from only one table at a time.** However, **the deletion may propagate to tuples in other relations if referential triggered actions are specified in the referential integrity constraints** of the DDL.

Depending on the number of tuples selected by the condition in the **WHERE** clause, zero, one, or **several tuples can be deleted by a single DELETE command.** **A missing WHERE** clause specifies that **all tuples in the relation are to be deleted;** however, the **table remains in the database as an empty table.** We must use the **DROP TABLE** command to remove the table definition.

```
U4A:  DELETE FROM    EMPLOYEE
      WHERE          Lname = 'Brown';
U4B:  DELETE FROM    EMPLOYEE
      WHERE          Ssn = '123456789';
U4C:  DELETE FROM    EMPLOYEE
      WHERE          Dno = 5;
U4D:  DELETE FROM    EMPLOYEE;
```

6.4.3 The UPDATE Command

The **UPDATE** command is used to **modify attribute values of one or more selected tuples.** As in the **DELETE** command, a **WHERE** clause in the **UPDATE** command **selects the tuples to be modified from a single relation.** However, **updating a primary key value may propagate to the foreign key values of tuples in other relations if such a referential triggered action is specified in the referential integrity constraints** of the DDL (see Section 6.2.2). An additional **SET** clause in the **UPDATE** command **specifies the attributes to be modified and their new values.**

```
U5:   UPDATE        PROJECT
      SET           Plocation = 'Bellaire', Dnum = 5
      WHERE        Pnumber = 10;
```

Several tuples can be modified with a single UPDATE command.

```
U6:  UPDATE  EMPLOYEE  
      SET    Salary = Salary * 1.1  
      WHERE  Dno = 5;
```

It is also possible to specify **NULL** or **DEFAULT** as the **new attribute value**. *Notice that each UPDATE command explicitly refers to a single relation only. To modify multiple relations, we must issue several UPDATE commands.*